

Toward the Frontier: Seventeen Measured Rungs Against a Frontier Canadian Fish Agent, and the Wall Where They Stopped

Allen Hsieh
FishAI project

September 2026

Abstract

The FishAI project’s previous cross-engine paper [2] measured its agent at 27.08% against SESTINA v1.0, the frontier bot of the FishLab project [3], played inside their C++ engine. This paper reports what a new agent line built to beat that number — Monet, seventeen rungs from v0.1 to the acceptance read — actually bought, under a protocol written before the first rung ran: every rung’s acceptance test written before it and, from the sixth rung, its predictions too; every number quoted per seed against a paired floor; and version 1.0 defined by an acceptance test and nothing else. **Monet v1.0 does not exist at the shipped vector.** The shipped bot reads 41.31% on twelve seeds and 40.11% on twelve others (14,400 games each), against a bar of 50.0%; it beats their v0.2 and v0.3 on every seed and loses to their v0.4, v0.5 and v0.6 on every seed; and a second adapter, written by an independent agent from FishLab’s two protocol documents alone, reproduces the win rate to the hundredth on twelve of twelve seeds and the game records byte for byte on ten. Three mechanisms shipped on the way: a calibrated per-card marginal with a joint chain over it and a licence term at low weight, a credit in the ask ranker for asks a hit-maximiser never plays, and a translation of the rule set’s declare compulsion that the bridge had been answering wrongly. Ten other rungs were measured and shipped nothing, each inside its floor; the belief axis closed on records studies rather than cells; the ask ranker closed on four fitted terms. The finding the read could not have predicted is where the gains went: against their v0.4, v0.5, v0.6 and SESTINA alike the ladder gained the same thirteen points, and none of the four is beaten — the wall at v0.4 stands where the first paper placed it. An attribution study locates the frontier’s margin — 40% of its extra sets in evenly dealt half-suits, 31% in steals from Monet’s majorities — and finds that at SESTINA’s own decisions Monet’s pick would hit 62.5% against SESTINA’s 55.7%: the hit chance is not the value of an ask. We report the ladder, the table, the instrument and its two disclosed defects, and the negative result as the deliverable.

1 Introduction

A laboratory that publishes a loss owes its readers the attempt. The first cross-engine measurement of this project [2] placed FishAI’s agent between FishLab’s v0.3 and v0.4 — a release four versions before search existed in that lineage — and left ~ 16 of the 22.92 points to even unattributed. The roadmap that followed [1] opened a new agent line, Monet, whose stated goal is one sentence: *Monet v1.0 beats SESTINA v1.0*. It then did something this project had not done before at that scale: it wrote the acceptance test for v1.0 before the first rung ran (six conditions, §4), it wrote a prediction for every rung before that rung’s cells were played, and it recorded every rung whether or not

anything shipped from it. Seventeen rungs later the acceptance test has been executed in full on the shipped bot, and this paper is its report.

The result is negative and the paper says so in the first line. The interest is in three other things. First, *what the fourteen points that were bought are made of* — three shipped mechanisms, one of which is not a policy change at all but a correction to how the bridge translates a rule the two engines state differently. Second, *why ten further rungs bought nothing*, which is the record of a belief programme measured to its end: the inference was studied on full-information records rather than cells, and every route into the belief that the records could price was priced at zero or below. Third, *where the ladder's gains went*. Against FishLab's non-searching v0.4, v0.5 and v0.6 and against SESTINA the ladder gained the same thirteen points, which says the shipped gains are worth the same against every opponent above the wall and bought no ground on the wall itself. The frontier's margin is elsewhere, and the attribution study of §5 says where.

Everything in this paper is a number in MONET.md [1], cited by section; the paper adds no measurement of its own. FishLab's repository carries no licence file, and nothing of theirs is reproduced here: their engine's game output is data, their documented guest-bot protocol is implemented, and their measured findings are cited as theirs.

2 The instrument

The bridge. FishAI's decision function is registered as a guest bot in FishLab's C++ engine over their JSON-line protocol, as in [2]. A cell is 200 deals $\times$ 6 seat rotations = 1,200 games on one seed; a read is twelve seeds, 14,400 games, quoted as a mean over seeds with its standard deviation and the paired standard error, against a paired floor of ± 2.00 points that the roadmap fixed before any rung (MONET.md §6.3). Every cell of the ladder is on the record with its seed, and every seed reported.

The controls. The mirror cell is not one of them — a policy against itself on duplicate deals is forced to 50% before a card is dealt, which [2] learned at the cost of a published wrong number. What is mandatory instead (MONET.md §6.2): a per-process count of every protocol op, cross-checked against the engine's own count of the guest's asks; a byte-exact null arm, the mechanism switched off, reproducing a known cell line for line before any new seed is spent; paired deals throughout; fault counters read by name; a calibration harness on every cell; and a completion check, because cells have died silently on handshake timeouts. From v0.14 every arm is also pinned *in-engine*: the recorded play is replayed through the policy at home and must agree with the bridge at 100.0% of the arm's decisions before its number is read.

The rule set, translated. us54 compels a cardless turn-holder to declare [5]; FishLab's engine offers it an optional poll instead, and its own ladder resolves what a guest declines. The bridge of [2] answered the compulsion into the optional poll, and the roadmap's records study of the declare (MONET.md §3.8f) found that every such answer below certainty spent a set to avoid a move that costs nothing. The translation — a compelled claim whose plan is not certain is answered *none*, and the host's ladder takes it — is worth +1.28 points paired on twelve seeds (SD 0.29, SE 0.08, ahead on all twelve). It is not a policy change and it is not credited to any rung; it is the bridge reading the rules correctly. Numbers measured before it and after it are on different bridges, and this paper never subtracts one from the other.

The records. FishLab’s engine keeps no game trace. A counters-only fork of their `match` binary, extended with one flag, appends the deal and the event trace of each game as a line of JSON after the game ends — ground truth read after the fact, changing no branch a policy can see, and controlled by requiring its whole cell output to equal the unmodified fork’s on the same cell (identical, at seed 90210). A loader at home reconstructs every seat’s view from the deal and the public log, checks every event against the engine’s own hand counts and every declare against its half-suit awards, and replays the policy at every ask decision of both sides: the counterfactual agrees with the bridge play at 100.0% of Monet’s 102,000 asks on the study’s records, and at 100% on both sides of a mirror cell in their engine. That agreement is what turns the records into an instrument: the same policy can be asked what it would have done at the *opponent’s* decisions, scored on the true deal.

The second arm. Condition 6 of the acceptance test requires the result to be reproduced by an independently built adapter on the same specification. An agent that had never seen this project’s adapter was given FishLab’s `BOT_PACKAGE.md` and `PLAY.md` and a one-page specification — the policy, the decision-seed rule, the compulsion translation — and built one. Checked at home it agreed with the engine’s own view constructor on 87,575 views with zero differences and on 16,449 of 16,449 decisions. Abroad it reproduces the first adapter’s win rate to the hundredth on all twelve seeds and its game records byte for byte on ten; the three games in 28,800 that differ are located in §4.

3 The ladder

Table 1 is the whole record, one row a rung, every number abroad against SESTINA v1.0 unless marked home. Three rungs shipped; the rest were measured and shipped nothing.

What shipped. The belief: a calibrated marginal over the unseen cards (`pModel: 'marginal'`), a chain that places a set’s open cards jointly over it when a declare is planned (`pAssignment: 'joint'`), and a licence term that conditions the hit probability on what a seat’s own asks have revealed, at a weight the 24-seed cell decided (`licenceLambda 0.3`; 0.6 read +1.88 over v0.4b on 22 of 24 seeds, and the pre-registered 0.3-against-0.6 confirmation read +1.19, SE 0.30, 19 of 24). Then the ask: a credit on the miss branch of the ask ranker (`contest: 0.6`) for asks into sets the opponents hold most of — the asks a hit-maximiser never plays — which is the one rung after v0.4c to clear the floor abroad, +4.04 paired on twelve pre-registered seeds and ahead on every one. Then the bridge’s translation of the compulsion. Everything else on the vector is v0.1’s Punter at **hard**.

What was measured and did not ship. Opponent reading through the prior (+0.47), communication (flat at home, never abroad), search at the affordable budget (+0.14 abroad; the correct posterior does not make a 100 ms search pay in either leaf), the determinized declare (+0.08), the exposure charge (+1.07, two standard errors above zero and two under the floor: real, and half a rung), the closing ask (+0.58, the same shape), the chase appetite (negative at every dose, closed on the fit without spending the confirmation seeds), and the calibrated licence likelihood (−0.74). Two more rungs were studies on the records that reached a verdict without a cell (§6). Each of the ten is on the record with the prediction written before it and the prediction’s score after.

The rule that did the work. A rung ships if its paired main effect clears +2.00 on twelve fresh seeds; inside the floor with its marker moved, it is recorded and not shipped. Four ranker terms were fitted in sequence on the same vector against a measured bound of +5.8 points (the counterfactual

Table 1: The Monet ladder against SESTINA v1.0. *Paired* is the arm minus its base over the same seeds, with the standard error over seeds; *ahead* counts seeds. The floor is ± 2.00 on twelve seeds, ± 2.83 on six. Bridge: *c* is the corrected bridge of [2]; *m* is that bridge with the compulsion translated (§2). A number on one bridge is never subtracted from a number on another. Sources: MONET.md §3.3–§3.8m.

rung	what it built	read	paired (SE)	ahead	bridge
v0.1	the fork: Bass v2.0’s Punter at hard , instrumented	27.08%	—	6 seeds	c
v0.3	licence conditioning of the hit probability, $\lambda = 0.60$	30.96%	+3.88	6 seeds	c
v0.4a	pCardAt, the calibrated per-card marginal (its own effect)	31.94%	+4.74	6 of 6	c
v0.4b	the joint chain over the marginal for the declare	32.75%	+0.81	—	c
v0.4c	λ back on the joint at 0.3, against 0.6 (24 seeds)	35.12%	+1.19 (0.30)	19 of 24	c
gate	the §3.5b panel on v0.4b, six seeds: the rule’s third row	32.75%	—	—	c
v0.5	opponent reading: ask-choice inference into the prior, $\kappa = 1$	34.92%	+0.47 (0.59)	7 of 12	c
v0.6	communication: asks chosen to reveal; the handoff (flat or behind at home)	home	—	—	home
v0.7	the search arm, 96 ms an ask, lock-only leaf	35.05%	+0.14 (0.67)	7 of 12	c
v0.8	the determinized declare	35.15%	+0.08 (0.08)	—	c
study	the attribution study, §5, twelve seeds	34.28%	—	—	c
v0.9	the contest credit on the ask ranker’s miss branch, 0.6	38.92%	+4.04 (0.47)	12 of 12	c
v0.10	the exposure charge on hits the opponents can take back	39.90%	+1.07 (0.54)	8 of 12	c
v0.11	the declare priced on records; the compulsion translated (the bridge, not the policy)	40.11%	+1.28 (0.08)	12 of 12	c→m
v0.12	the closing ask: a credit for bringing a held set within reach	40.69%	+0.58 (0.23)	9 of 12	m
v0.13	the chase appetite: every dose loses on the fit	fit	−1.08 to −7.56	none	m
v0.14	the assignment, a records study: the axis closed on calibration, the channel’s ceiling 0.33 sets a game	records	—	—	m
v0.15	the deduction fix, on records: exact inference reads worse	records	—	—	m
v0.16	the licence likelihood, calibrated on records	40.58%	−0.74 (0.46)	4 of 12	m
v0.17	the acceptance read on v0.9’s vector, against a bar of 50.0	41.31%	—	0 of 24	m

surplus of §5); the first took +4.04 of it, the next three +1.07, +0.58 and less than zero. The roadmap’s own reading is that the ask ranker was finished by v0.13 and that the last term was the informative one: it went looking for the missing chases, found that a seat cannot identify them from its view, and paid a hundred points a decision to learn it (MONET.md §3.8i).

4 The acceptance test

Monet v1.0 exists when, and only when (MONET.md §3.9): (1) its win rate against SESTINA v1.0 is at least 50.0% pooled over twelve seeds of 200 deals at the ± 2.00 floor — never on three seeds, never on one cell; (2) every seed is reported and the SD across them published; (3) the panel is monotone and Monet beats FishLab’s v0.2 through v0.6 as well, because a bot that beats SESTINA and loses to v0.5 has been fitted, not built; (4) declare accuracy is at least 98.0% and every fault counter zero; (5) every control of §2 passes, the mirror not among them; (6) the result is reproduced by a second, independently built arm. The read was pre-registered on 2026-09-06 with six predictions and executed the same day (MONET.md §3.8m).

Table 2: The acceptance table, read on v0.9’s vector. Sources: MONET.md §3.8m.

condition	reads	met
1 · $\geq 50.0\%$ vs SESTINA v1.0, 12 seeds	41.31% on one twelve (SD 1.29, range 39.42–43.00); 40.11% on another (SD 1.55, range 37.25–42.33)	no , by 8.7 and 9.9
2 · every seed reported, SD published	24 seeds on record	yes
3 · the panel monotone, v0.2–v0.6 beaten	monotone; v02 74.18 and v03 70.40 beaten on every seed; v04 47.62, v05 46.81, v06 44.72 beaten on none	no , three of five
4 · declare accuracy $\geq 98\%$; counters zero	99.20–99.82% on 73 cells; the fourteen counters zero on every cell of the restated adapter	yes
5 · the controls	the per-process ask count equal to the engine’s on 73 of 73 cells; the null arm identical; calibration inside its bars on 12 of 12; every cell complete	yes
6 · a second, independent arm	the same win rate to the hundredth on 12 of 12; records identical on 10 of 12; 3 games of 28,800 differ at one located position, no result changed	yes

Table 3: The panel on the read’s twelve seeds: Monet v0.9 against FishLab’s lineage in their engine. Means over seeds; the inherited row is Monet v0.1’s level on the corrected bridge, from [2] and MONET.md §6.4, and the gain is a difference of levels across the bridge translation, quoted as the roadmap quotes it and not as a paired contrast. Source: MONET.md §3.8m.

	v02	v03	v04	v05	v06	SESTINA
Monet v0.9, mean	74.18	70.40	47.62	46.81	44.72	41.31
SD over seeds	1.15	1.07	1.86	1.42	1.39	1.29
seeds above 50%	12 of 12	12 of 12	0 of 12	0 of 12	0 of 12	0 of 12
sets a game, Monet – opponent	5.41 – 3.59	5.22 – 3.78	4.44 – 4.56	4.40 – 4.60	4.32 – 4.69	4.19 – 4.81
Monet v0.1, inherited	67.42	62.06	34.25	33.31	32.86	27.83
the ladder’s gain	+6.8	+8.3	+13.4	+13.5	+11.9	+13.5

Conditions 1 and 2. Two reads of the same vector on seed sets that share nothing: 41.31% and 40.11%, 1.20 points apart with a standard error of 0.58 between them — two seeds’ worth of the same bot, not two bots. Every one of the twenty-four seeds is under 44%.

Condition 3: the panel, and the wall. Sixty cells, 72,000 games, on the same twelve seeds as the read (Table 3). The order is monotone and the panel is what the roadmap said it would be from its first panel at v0.4b: two rungs Monet clears, then a wall at their v0.4 — *“the rung that carries the lineage’s whole strength”* [1] — with the frontier six points behind the wall. The number the read could not have predicted is in the last row. Against v0.1’s inherited levels the ladder gained +6.8 and +8.3 against the two lineage bots it already beat, and +13.4, +13.5, +11.9 and +13.5 against v0.4, v0.5, v0.6 and SESTINA: the same thirteen points on each of the four it does not beat. The shipped gains — the licence weight, the contest credit, the compulsion translated — are worth the same against the wall as against the frontier, and bought no ground on the wall itself. From v0.4 to SESTINA Monet loses between 0.1 and 0.6 sets a game.

Condition 4. Declare accuracy on the 73 cells of the read runs 99.20–99.82% against a bar of 98.0%; the parity with the frontier that [2] found still holds. The fourteen fault counters are zero on every cell played by the adapter as restated for this read (§7), and the engine’s own **limit hits** line reads 0% everywhere.

Condition 5. The per-process count of the guest’s asks equals the engine’s own count of them on 73 of 73 cells — the cross-check the controls ask for, exact — and the pass op is exercised on every cell, which is the tripwire that would have caught the defect of [2] in an hour. The null arm, replayed on the identity seed before any new cell, reproduced the base cell line for line. Calibration, the guest’s own asks believed against realised: aggregate bias 0.0046–0.0096 against a bar of 0.01 on twelve of twelve cells, worst decile 0.065–0.098 against 0.10 on twelve of twelve — with the same shape the roadmap first measured at v0.3, the belief a little too sure of its long shots and not sure enough of its good ones.

Condition 6, and the three games. The independent adapter’s twelve cells match the first adapter’s win rate on every seed and its records byte for byte on ten. On the other two, three games of 28,800 differ, each at one event, and each at the same position: a seat has just declared its last cards, holds the turn with a pass pending, and the host polls the table. The first adapter’s cardless seat passes and the next teammate polled declares the team’s next certain set; the second arm’s seat declares it itself before passing. Same set, same side, same final score. **us54**’s own reducer closes the declare window for the duration of a pending pass [5], so this is a position FishAI’s rules never reach and the specification did not name; the two authors resolved it differently, and the two resolutions differ on three events in 28,800 games and on no result. The condition is met, and the position is now specified.

5 Where the sets go

Before v0.9 the roadmap ran a study rather than a rung (MONET.md §3.8c): full-information records of Monet v0.4c against SESTINA on twelve seeds, split four ways by one instrument, with the next rung to be chosen by a rule written before the tables were read. Monet won 34.28% of those games and 3.22 sets a game to SESTINA’s 4.25, a gap of 1.033 sets (SE 0.021).

By the deal. Split each half-suit by how many of its six cards Monet’s side was dealt, the set gap is +0.608 a game on Monet’s majorities (four to six cards), −0.414 on the evenly dealt threes, and −1.227 on SESTINA’s majorities. Of SESTINA’s 1.09 extra *cached* sets a game, the even half-suits carry 0.43 (40%; a 3–3 deal is won by SESTINA 58 to 42), SESTINA’s steals from Monet’s majorities 0.34 (31%), and its better conversion of its own majorities 0.31 (29%). Gifts — sets handed over by a wrong declare — run the other way and are small. At the clinch Monet holds 0.30 sets a game locked and uncashed to SESTINA’s 0.10, and holds a lock 7.84 events to SESTINA’s 4.34: the game ends while Monet is still collecting, which is the proof-latency finding of [2] again, now bounded at about 0.2 sets a game.

The counterfactual, and the third case. At every ask decision of SESTINA the records were replayed through Monet’s policy, from SESTINA’s hand and the public log, and scored on the true deal beside the ask SESTINA actually made. The study’s pre-registration named three cases and the tables returned the third, decisively: Monet’s pick would hit 62.5% where SESTINA’s asks hit 55.7% — a surplus of 6.81 points, SE 0.08 over twelve seeds, the same on every seed and in every holding bucket, and present in SESTINA’s own mirror (54.7 against 49.4). SESTINA agrees with Monet’s pick 42.8% of the time, and where the two differ SESTINA’s ask hits 41.1% to the pick’s 53.0%. SESTINA leaves a quarter of the available hit chance on the table at every decision and wins two games in three. **The hit chance is not the value of an ask.** What SESTINA’s asks buy shows in the same tables: its positions are richer — the same Monet policy hits 52.4% at Monet’s

own positions and 62.5% at SESTINA's — and they get richer as the game goes, where Monet's are flat. A hit publishes a card's location to the whole table and a miss publishes a licence; SESTINA prices both, and Monet's ranker priced neither.

What v0.9 bought, and what it paid. The contest credit moves 15.2% of Monet's asks, at the decisions where the greedy pick is poorest: there v0.4c's pick would hit 22.4% and the priced ask hits 24.4% — not worse asks, but asks placed elsewhere, into sets where Monet holds one to three cards and away from the five-card closing asks that on a miss hand the turn over with the set exposed. Monet keeps the turn longer (41.7 asks a game from 39.2), the even half-suits go to Monet 46.1% of the time from 42.0%, its four-card majorities convert 62.1% from 57.2%, and the lock hold falls from 7.90 events to 5.93 because the run that formed the lock is still going. It also pays what the study said a hit costs, and more: its hits are taken back 43.4% of the time from 41.0%, and its misses are made at more dangerous targets. Two of the five markers the pre-registration wrote came in as written; the mechanism that won was the contest, not the take-back the study had pointed at, and the record says so. FishLab's own finding about SESTINA — that ask accuracy is anti-correlated with strength near the optimum [3] — is confirmed in kind: Monet's accuracy *rose* under the credit, because the decisions it changed were the poorest ones.

6 Why the belief axis closed

The ladder's last three rungs, v0.14 to v0.16, were about the belief — what Monet thinks each seat holds — and the first two of them never needed a cell. With the records of §5 and a seam that runs the real policy on an injected belief with nothing in the engine touched, a candidate belief can be scored against the true deal at every decision Monet made, and a change that does not move the belief where the belief is wrong is dead before it plays a game.

The assignment (v0.14). Is the side assignment — which team holds a missing card — the binding constraint? Studied on two twelve-seed corpora of records (MONET.md §3.8j): the assignment is calibrated, the exact pricer that must first reproduce known answers fails to, and the ceiling of the whole channel under a *perfect* belief is 0.33 sets a game, half a point of win rate at the class prior. The axis closed on its calibration, and the candidate that survived was a correctness fix to the deduction.

The deduction (v0.15). The fix was studied before any code: exact conditioning on the model's own constraints, in place of the approximation the engine ships. It reads *worse* on the records — the pooled Brier score rises rather than falls — and the pre-registration was wrong about the direction of nearly everything it predicted (MONET.md §3.8k). The model, not the inference, is the error: a seat that has asked into a set holds about 1.5 of its alive cards whatever their number, and the model's licence carries no such likelihood.

The licence likelihood (v0.16). So the likelihood was calibrated on the records — three knob families through the same seam, fit on three seeds and held out on twenty-one — and the one that survived its falsifiers repaired the belief exactly where v0.15 had located the error. Abroad it lost: 40.58% against the base's 41.31%, -0.74 paired, inside the floor, behind on sets on ten of twelve seeds with its wrong declares doubled (0.0125 to 0.0239 a game, worse on every seed, seven standard errors). The loss is located and it is a lesson about instruments rather than beliefs. The declare places a set's open cards by a chain over the same table, and that path was never among the falsifiers,

which all read pooled calibration; on the same declare windows the repaired table prices one-guess plans above 0.9 nineteen times as often as the shipped one, and they come right 43% of the time against the shipped table’s 70%. **A pooled Brier score cannot see a selected population.** The declare gate selects exactly the cells where a table is over-confident, and no belief change on this ladder may now ship without the plan’s reliability read through the planner itself on the same windows.

The count. Seven consecutive measured negatives, v0.10 through v0.16, the last with the belief repaired where the error was and worth nothing in play, closed the axis (MONET.md §3.81). An oracle arm that shares its team’s true hands and cashes every lock the instant it exists bounded every belief mechanism at 38.28% on the v0.4-era ask policy (MONET.md §3.4a); v0.9 reads above it because the ask policy changed, not because the bound was wrong.

7 Method notes

Pre-registration, scored. Every rung’s section in [1] carries a pre-registration written before its cells and a record that scores the predictions after. The scores are reported whether or not they flatter the rung: v0.9’s mechanism was not the one its markers pointed at (two of five as written); v0.15’s pre-registration missed on the sign of its headline; v0.17’s six predictions hit five and missed the sixth by a tenth of a point on one cell. The acceptance read predicted that its own condition 6 would be met on the number and might differ in the letter at a declare-window position, and it did, at a pending pass.

A stale instrument, disclosed. The bridge adapter carries a cross-check that restates each declare’s card placement locally and counts a mismatch with the engine’s own. It had restated by the capacity-greedy rule since v0.1, and the engine has placed cards by the joint chain since v0.4b, so the counter read non-zero on every kept log from v0.10 on — 220 to 459 a cell against about 11,000 declares — and was read as noise. It was a stale instrument, not a bot fault; the wire always carried the engine’s own placement. For the acceptance read the check was restated through the engine’s planner, the greedy comparison kept as information, and the restated adapter replayed the identity cell byte-identical to the original: on that cell the information counter reads 360 and 723, exactly the stale counts. The disclosure is in the record, and every record from here reports the fourteen counters by name.

Where the counts come from. The op counts are summed from per-process files, never from the shared log, which loses about one process in thirty-six; the read’s own pre-registration had quoted bands from the log and its record restates them from the files, where the ask count equals the engine’s on every cell. The controls of [2] were written from exactly this failure.

Cost. Monet’s policy decides an ask in about 0.14 ms and a six-seat game in about 82 ms at home. The search arm at its pre-registered form costs 96 ms an ask decision against a budget of 100; its wider form twice the budget and its deeper form 1.7 times, and both were out on the cost rule whatever they read (−0.08 and +0.60 sets a pair at home, SE about 0.34).

8 What is left

The residual is position. The decomposition of [2] — Monet picks better moves than the frontier at matched positions (+5.49) and arrives at worse ones (−10.38) — was measured on the defective bridge and has not been re-recorded, and the roadmap labels it a hypothesis. Everything this ladder measured is consistent with it. The belief is calibrated and repairing it buys nothing; the ask ranker is finished at four terms; a hit-maximising policy at the frontier’s own decisions would out-hit the frontier by seven points and the frontier wins anyway. What the frontier buys with its asks is a richer position three moves later, and that is not a knob on a one-step ranker.

Search at the affordable budget is a null. The one channel with a ceiling above the gap is lookahead, and at 100 ms an ask the correct posterior does not make it pay (+0.14 abroad). At that horizon the set differential almost never moves — nine candidate means in ten read exactly zero — and a locked set counted as a set does not change what the ranker’s own alternatives are worth. What a search that pays would need is a leaf that prices what an ask publishes, which is the same object §5 found the frontier pricing and the ranker not.

The even half-suits. 40% of the frontier’s margin sits in the 3–3 deals, which v0.9 moved from 42.0% to 46.1% and no rung has attacked as a target. It remains the leading candidate for a rung, at the prior of a ranker term, which on this record is under the floor.

The decision. The roadmap wrote the fork for this outcome before the first rung: ship the fast policy in the browser at its honest maximum and publish the negative result, or fund a lab-only search arm beyond the affordable budget. The owner took the first on 2026-09-06: v0.9 is what the site’s table seats, this paper is the artefact, and the direction from here is that a rung is chosen by what it is measured to add to the number, with the same protocol — pre-registered, paired, floored, reproduced — deciding what that is.

9 Threats to validity

One opponent. The target is one frozen specification of one project’s bot; nothing here is a claim about Canadian Fish in general, and the panel exists to show that the gains are not a fit to that one opponent — they are the same size against four. *One rule set.* Every number is conditional on **us54** as translated into a host that plays nine half-suits to the end where **us54** clinches at five; every figure here is read over the pre-clinch record, and the host’s **win-condition** warnings on every cell are the expected consequence of that difference. *Two bridges.* The compulsion translation moved every Monet number by about +1.3 points; Table 1 labels each row’s bridge and the paper subtracts across the seam nowhere except in the panel’s inherited row, which the roadmap quotes as a difference of levels. *The second arm is a second adapter, not a second engine.* Condition 6 reproduces the bridge, which is what it was written to do; the engine is FishLab’s on both sides. *The position decomposition is on the defective bridge* and is a hypothesis, labelled as one. *Statistical resolution.* Twelve seeds resolve ± 2.00 paired; two of the rungs that shipped nothing read between one and two points with standard errors of a quarter to half a point, which is a real effect under the floor, and the record says so rather than rounding either way.

Data and code

The engine, the policy, the records instrument (`scripts/attribute.mjs`), the tests that pin every shipped vector and its forward bank, and MONET.md with every cell of the ladder are in the FishAI repository. The bridge — FishLab’s engine built from source, the recording fork, the arms and the game records — is outside it, because FishLab’s repository carries no licence file; it is archived by the owner outside the public tree and every number here can be re-read from it.

References

- [1] FishAI project. MONET.md — the roadmap from Monet v0.1 to v1.0. Repository document, 2026. §0 the target and the fork; §3.3–§3.4 the belief rungs; §3.5b the gate; §3.8a the search arm; §3.8c the attribution study; §3.8d v0.9; §3.8e–§3.8i the ranker rungs and the bridge translation; §3.8j–§3.8l the belief studies and the declare pin; §3.8m the acceptance read; §3.9 the definition of v1.0; §6.2 the controls; §8.3 the decisions.
- [2] A. Hsieh. Against the frontier: a cross-engine measurement and the bridge defect that nearly buried it. FishAI project, 2026. The 27.08% baseline, the mirror-cell lesson, the oracle ceiling and the position decomposition this paper inherits.
- [3] dylann4500. FishLab — a Canadian Fish engine, bot family and evaluation lab. GitHub repository, <https://github.com/dylann4500/FishLab>. The C++ engine and its identity controls, the SESTINA v1.0 frozen specification, the v0.2–v0.6 lineage, the guest-bot package protocol (`docs/BOT_PACKAGE.md`, `docs/PLAY.md`), and the finding on ask accuracy near the optimum. The repository carries no licence file; nothing from it is reproduced here, and its measured findings are cited as theirs.
- [4] FishAI project. WHY-FISHAI-LOSES.md — the account of the cross-play deficit. Repository document, 2026. The proof-latency channel and the lineage placement the ladder started from.
- [5] FishAI project. RULES_US54.md — the “US student” 54-card variant. Repository rules specification with derived consequences and test vectors, 2026. §3.1 the pending pass, §3.2 the declare compulsion, §4 running out of cards.
- [6] A. Hsieh. FishAI v2.0: the three-sided ask. FishAI project, 2026. The Bass v2.0 policy Monet v0.1 forked, and the defusal term every roster style carries.